

Reinforcement Learning

For the second part of this project we will be implementing a simple Q-Learning algorithm on an RL environment called Cart Pole. The idea of Q-Learning is to try to estimate the expected future reward or Q-value of taking a certain action. Then at any given step we take the action with the most expected future reward.

In reinforcement learning, we refer to algorithms that attempt to solve environments as "agents", so in this part of the project we will be making a Deep Q Network Agent that will solve the Cart Pole environment.

In [1]:

```
!pip install gym tqdm

Collecting gym
  Downloading
https://files.pythonhosted.org/packages/9b/50/ed4a03d2be47ffd043be2ee514f3295d98a30fe2d1b9c61dea5a9d861/gym-0.10.5.tar.gz (1.5MB)
  100% |████████████████████████████████████████| 1.5MB 566kB/s
Collecting tqdm
  Downloading
https://files.pythonhosted.org/packages/78/bc/de067ab2d700b91717dc5459d86a182df31abfb90ab01a5a5a5ce30b4/tqdm-4.23.0-py2.py3-none-any.whl (42kB)
  100% |████████████████████████████████████████| 51kB 1.6MB/s
Requirement already satisfied: numpy>=1.10.4 in
/anaconda3/envs/py36/lib/python3.6/site-packages (from gym)
Requirement already satisfied: requests>=2.0 in
/anaconda3/envs/py36/lib/python3.6/site-packages (from gym)
Requirement already satisfied: six in
/anaconda3/envs/py36/lib/python3.6/site-packages (from gym)
Collecting pygame>=1.2.0 (from gym)
  Downloading
https://files.pythonhosted.org/packages/1c/fc/dad5eaaab68f0c21e2f906a94ddb98662cc5a654eee404d59554ce0fa/pyglet-1.3.2-py2.py3-none-any.whl (1.0MB)
  100% |████████████████████████████████████████| 1.0MB 827kB/s
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in
/anaconda3/envs/py36/lib/python3.6/site-packages (from requests>=2.0->gym)
Requirement already satisfied: idna<2.7,>=2.5 in
/anaconda3/envs/py36/lib/python3.6/site-packages (from requests>=2.0->gym)
Requirement already satisfied: urllib3<1.23,>=1.21.1 in
/anaconda3/envs/py36/lib/python3.6/site-packages (from requests>=2.0->gym)
Requirement already satisfied: certifi>=2017.4.17 in
/anaconda3/envs/py36/lib/python3.6/site-packages (from requests>=2.0->gym)
Collecting future (from pygame>=1.2.0->gym)
  Downloading
https://files.pythonhosted.org/packages/00/2b/8d082ddfed935f3608cc61140df6de4ealbc3ab52fb6c29ae3e81e85/future-0.16.0.tar.gz (824kB)
  100% |████████████████████████████████████████| 829kB 1.0MB/s
Building wheels for collected packages: gym, future
  Running setup.py bdist_wheel for gym ... done
  Stored in directory:
/Users/chapatel/Library/Caches/pip/wheels/cb/14/71/f4ab006b1e6ff75c2b54985c/
```

```
d0644fffe9c1dddc670925
```

```
Running setup.py bdist_wheel for future ... done
```

```
Stored in directory:
```

```
/Users/chapatel/Library/Caches/pip/wheels/bf/c9/a3/c538d90ef17cf7823fa51fc70a7a910a80f6a405bf15b1a
```

```
Successfully built gym future
```

```
Installing collected packages: future, pygame, gym, tqdm
```

```
Successfully installed future-0.16.0 gym-0.10.5 pygame-1.3.2 tqdm-4.23.0
```

```
You are using pip version 9.0.1, however version 10.0.1 is available.
```

```
You should consider upgrading via the 'pip install --upgrade pip' command.
```



Part 1: Setup the Environment

In [2]:

```
import gym
env = gym.make('CartPole-v0')
```

WARN: gym.spaces.Box autodetected dtype as <class 'numpy.float32'>. Please provide explicit dtype.

Part 2: Create The DQN Agent

In [3]:

```
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Activation
from collections import deque
import random
from keras.optimizers import Adam

import numpy as np

class DQNAgent:

    def __init__(self, env, replay_size=1000, epsilon=1.0, epsilon_min=0.01,
                 epsilon_decay=0.995, gamma=0.99):
        self.state_size = env.observation_space.shape[0]
        self.num_actions = env.action_space.n
        self.model = self.build_model()
        self.replay_buffer = deque(maxlen=replay_size)
        self.epsilon = epsilon
        self.epsilon_min = epsilon_min
        self.epsilon_decay = epsilon_decay
        self.gamma = gamma

    def build_model(self):
        model = Sequential()
        # TODO: add 2 dense layers each with 32 neurons, the input dim to the first
        # layer should be the state size, also add relu activations, for both these layers
```

```

        # Then add another Dense layer with num_actions neurons.
        # Then use model.compile to compile the model with mse loss and an
Adam optimizer
        # with learning rate 0.001.
        model.add(Dense(32, input_dim = self.state_size))
        model.add(Activation("relu"))
        model.add(Dense(32))
        model.add(Activation("relu"))
        model.add(Dense(self.num_actions))
        keras.optimizers.Adam(lr=0.001)
        model.compile(optimizer = "Adam", loss="mse")

    return model

    def action(self, state):
        # Whenever a random number between 0 and 1 is less than epsilon we
want to return
        # a random action. This means that with probability epsilon we retu
rn a random action.
        if np.random.random() <= self.epsilon:
            return np.random.randint(self.num_actions)
            #TODO: return random action here
        # Now we want to use our model to get the q values
        # HINT: we want to do prediction

        q_values = self.model.predict(state)
        return np.argmax(q_values[0])

    def add_to_replay_buffer(self, state, action, reward, next_state, done)
:
        self.replay_buffer.append((state, action, reward, next_state,
done))

    def train_batch_from_replay(self, batch_size):
        # if we don't have enough samples in our replay buffer just return
        if len(self.replay_buffer) < batch_size:
            return False
        # TODO: randomly sample batch_size samples from the replay buffer
        # hint: use random.sample
        minibatch = random.sample(self.replay_buffer, batch_size)
        for state, action, reward, next_state, done in minibatch:
            target = reward
            if not done:
                next_Qs = self.model.predict(next_state)[0]
                # TODO: we want to add to our target GAMMA * max
Q(next_state)
                target += self.gamma * np.max(next_Qs)

            # our target should only take into account the current action
            # so we set all the Q values except the current action, to the
            # current output of our model so that they get ignored in the l
oss function.
            target_Qs = self.model.predict(state)
            target_Qs[0][action] = target
            self.model.fit(state, target_Qs, epochs=1, verbose=0)

        # Now we want to slowly decay how many random actions we take
        # to do this we can multiply epsilon by our epsilon decay parameter
        # each iteration

```

```
if self.epsilon > self.epsilon_min:
    self.epsilon *= self.epsilon_decay
```

Using TensorFlow backend.

Part 3: Train the Model

In [4]:

```
agent = DQNAgent(env)
```

In [5]:

```
from tqdm import tqdm

done = False
batch_size = 32
num_episodes = 800

for episode in tqdm(range(num_episodes)):
    state = env.reset()
    state = np.reshape(state, [1, agent.state_size])

    for t in range(200):
        action = agent.action(state)
        next_state, reward, done, _ = env.step(action)
        reward = reward if not done else 100
        next_state = np.reshape(next_state, [1, agent.state_size])
        agent.add_to_replay_buffer(state, action, reward, next_state, done)
        # TODO: add this sample to the replay buffer

        state = next_state

        # TODO: train on a batch from the replay buffer
        agent.train_batch_from_replay(batch_size)
    if done:
        break
```

100%|██████████| 800/800 [34:59<00:00, 2.62s/it]

Part 4: Test the Model

In [6]:

```
#TODO: set the agent's epsilon so that we don't take any random actions.
for _ in range(10):
    state = env.reset()
    state = np.reshape(state, [1, agent.state_size])
    agent.epsilon = -1
    total_reward = 0
    for t in range(200):
        action = agent.action(state)
        next_state, reward, done, _ = env.step(action)
        total_reward += reward
        state = np.reshape(next_state, [1, agent.state_size])
        # TODO: if you want to see the rendered version of your agent running
```

```
# uncomment this line
#env.render()
print(total_reward)
```

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

40.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

23.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

27.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

21.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

35.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

21.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

24.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

46.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

20.0

WARN: You are calling 'step()' even though this environment has already returned done = True. You should always call 'reset()' once you receive 'done = True' -- any further steps are undefined behavior.

27.0

Part 5: Writeup

Now for the writeup portion write a paragraph of your understanding of how Deep Q Learning works.

Q-learning uses a simple update rule to perform q-value iteration, which allows us to bypass the need to keep track of values, transition functions, and reward functions. We use Deep Q-Learning to approximate our Q-value function with the use of a Neural Network. We choose the neuron from our network that has the highest value and take an action corresponding to this neuron.

STYLE TRANSFER

In this project we are going to be using Convolutional Neural Networks to implement Neural Style Transfer, a technique for creating a new image with the content of one input image and the style of another input. The idea behind style transfer is as follows: take three input images, one our style image, one our content image, and one output image which starts as random noise and we iteratively update it until it looks like the content of the content image in the style of the style image. To do this we run all three images through a pretrained VGG16 model trained to classify images. Then for a selected convolutional or pooling layer of the VGG16 model we compare the activations (values of the neurons) at that layer for the three different images. Specifically, we have what is called a feature reconstruction loss that compares the activations of the current output image and the content image, and what is called a style loss that compares the activations of the current output image and the desired style image. Then we use the gradient of these loss functions to update our current output image. Hopefully, that gives you an overview of what we will be doing, and you should gain a more in depth understanding as we go along.

The paper we will be implementing is found here:

<https://arxiv.org/pdf/1508.06576.pdf>

Part 1: Build Model and Define Losses

First we want to initialize a VGG16 model we can use for style transfer.

In [28]:

```
from keras.applications import vgg16
from keras.layers import Input, Concatenate
from keras.models import Model, Sequential
from keras import backend as K
import tensorflow as tf
```

In [29]:

```
K.clear_session()
content_input = Input(batch_shape=(1, 224, 224, 3))
style_input = Input(batch_shape=(1, 224, 224, 3))
output_tensor = tf.get_variable("output_tensor", [1, 224, 224, 3])
output_input = Input(tensor=output_tensor)
## TODO: use a concatenate layer to concatenate the three inputs on the
first axis.
input_tensor = tf.concat([style_input, output_input, content_input], 1)
```

If you get an error for the cell below about SSL PROTOCOL VERSIONS or something similar you can try downloading this file <https://github.com/fchollet/deep-learning->

[models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5](#) and putting it in the folder `~/keras/models` on your computer. Then the cell below should work after that

In [30]:

```
# We now create a pretrained VGG 16 model, which is really easy to do in Keras
# include_top=False ensures we don't use the fully connected layers.
vgg_model = vgg16.VGG16(input_tensor=input_tensor, weights='imagenet', include_top=False)
# We can now look at the structure of this model
vgg_model.summary()
```

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	(None, None, None, 3)	0
<hr/>		
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
<hr/>		
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
<hr/>		
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
<hr/>		
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
<hr/>		
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
<hr/>		
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
<hr/>		
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
<hr/>		
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

In [31]:

```
print([layer.name for layer in vgg_model.layers])
```

```
['input_4', 'block1_conv1', 'block1_conv2', 'block1_pool', 'block2_conv1',  
'block2_conv2', 'block2_pool', 'block3_conv1', 'block3_conv2',  
'block3_conv3', 'block3_pool', 'block4_conv1', 'block4_conv2',  
'block4_conv3', 'block4_pool', 'block5_conv1', 'block5_conv2',  
'block5_conv3', 'block5_pool']
```

In [115]:

```
# now select one of the above listed layers to be the layer to use for  
content information  
# and select some number of layers (maybe 2 or 3 layers) from the above  
layers to be the  
# style information. If you choose layers closer to the input this will use  
  
# more simplistic features, and choosing layers closer to the end will use  
# more complicated  
# abstracted features.  
content_layer = vgg_model.layers[4].name  
style_layers = [  
    vgg_model.layers[len(vgg_model.layers) - 3].name,  
    vgg_model.layers[len(vgg_model.layers) - 2].name,  
    vgg_model.layers[len(vgg_model.layers) - 1].name  
]  
# you can also play with the_content and style loss weights if you want  
to. this will effect  
# how stylized vs similar to the content image the output will look.  
content_loss_weight = 100  
style_loss_weight = 750  
content_layer, style_layers
```

Out[115]:

```
('block2_conv1', ['block5_conv2', 'block5_conv3', 'block5_pool'])
```

In [116]:

```
layers_dict= dict([(layer.name, layer.output) for layer in vgg_model.layers  
)  
layers_dict
```

Out[116]:

```
{'block1_conv1': <tf.Tensor 'block1_conv1/Relu:0' shape=(1, 672, 224, 64) d  
type=float32>,  
 'block1_conv2': <tf.Tensor 'block1_conv2/Relu:0' shape=(1, 672, 224, 64) d  
type=float32>,  
 'block1_pool': <tf.Tensor 'block1_pool/MaxPool:0' shape=(1, 336, 112, 64)  
dtype=float32>,  
 'block2_conv1': <tf.Tensor 'block2_conv1/Relu:0' shape=(1, 336, 112, 128)  
dtype=float32>,  
 'block2_conv2': <tf.Tensor 'block2_conv2/Relu:0' shape=(1, 336, 112, 128)  
dtype=float32>,  
 'block2_pool': <tf.Tensor 'block2_pool/MaxPool:0' shape=(1, 168, 56, 128)  
dtype=float32>,  
 'block3_conv1': <tf.Tensor 'block3_conv1/Relu:0' shape=(1, 168, 56, 256) d  
type=float32>,  
 'block3_conv2': <tf.Tensor 'block3_conv2/Relu:0' shape=(1, 168, 56, 256) d  
type=float32>,  
 'block3_conv3': <tf.Tensor 'block3_conv3/Relu:0' shape=(1, 168, 56, 256) d
```



```

type=float32>,
'block3_pool': <tf.Tensor 'block3_pool/MaxPool:0' shape=(1, 84, 28, 256) dtype=float32>,
'block4_conv1': <tf.Tensor 'block4_conv1/Relu:0' shape=(1, 84, 28, 512) dtype=float32>,
'block4_conv2': <tf.Tensor 'block4_conv2/Relu:0' shape=(1, 84, 28, 512) dtype=float32>,
'block4_conv3': <tf.Tensor 'block4_conv3/Relu:0' shape=(1, 84, 28, 512) dtype=float32>,
'block4_pool': <tf.Tensor 'block4_pool/MaxPool:0' shape=(1, 42, 14, 512) dtype=float32>,
'block5_conv1': <tf.Tensor 'block5_conv1/Relu:0' shape=(1, 42, 14, 512) dtype=float32>,
'block5_conv2': <tf.Tensor 'block5_conv2/Relu:0' shape=(1, 42, 14, 512) dtype=float32>,
'block5_conv3': <tf.Tensor 'block5_conv3/Relu:0' shape=(1, 42, 14, 512) dtype=float32>,
'block5_pool': <tf.Tensor 'block5_pool/MaxPool:0' shape=(1, 21, 7, 512) dtype=float32>,


```

Loss Functions

We want to define our style transfer losses now. First, we are going to define a feature reconstruction loss based on our content features and our output features. Using tensorflow functions implement the following loss function:

$$\frac{1}{2} \sum_{i,j,k} (F_{ijk} - P_{ijk})^2$$

where F is the 3D tensor of content features and P is the 3D tensor of our output image features.

HINT: `tf.reduce_sum` and `tf.square` will be helpful here.

In [117]:

```

def feature_reconstruction_loss(content_img_features, output_img_features):
    """Takes a tensor representing a layer of VGG features from the content
    image
    and a tensor representing a layer of VGG features from the current output
    image and returns a loss value.
    """
    return 0.5 *
    tf.reduce_sum(tf.square(content_img_features-output_img_features))

```

Now we wish to define our style loss function. First, we have to take our features and represent them as a Gram Matrix, for more information on Gram Matrices and this loss function you can read the paper if you like. Then we wish to implement the loss function:

$$\frac{1}{4H^2W^2C^2} \sum_{ij} (G_{ij} - A_{ij})^2$$

where G is the Gram matrix of the output image features and A is the Gram Matrix of the style image features. Note that we have written a Gram matrix function for you so you only need to call it.

In [118]:

```
def gram_matrix(x):
    # make channels first dimension
    x = tf.transpose(x, (2, 0, 1))
    # flatten everything but channels so x is now (C, H*W)
    x = tf.reshape(x, tf.stack([-1, tf.reduce_prod(tf.shape(x) [1:])]))
    return tf.matmul(x, tf.transpose(x))
```

In [119]:

```
def style_loss(style_img_features, output_img_features, img_shape):
    """Takes a tensor representing a layer of VGG features from the style i
    mage and a tensor
    representing a layer of VGG features from the current output image and
    returns
    the style loss for these features.
    """
    G = gram_matrix(output_img_features)
    A = gram_matrix(style_img_features)
    return tf.reduce_sum(tf.square(G-A))/(4* img_shape[0]**2 *img_shape[1]**
2 *img_shape[2]**2)
```

In [120]:

```
print(layers_dict[content_layer].shape)
print(tf.transpose(layers_dict[content_layer])[2, :, :, :])

(1, 336, 112, 128)
Tensor("strided_slice_30:0", shape=(112, 336, 1), dtype=float32)
```

In [121]:

```
content_features = layers_dict[content_layer]
content_img_features = tf.transpose(content_features)[0, :, :, :]
# output_content_features = content_features[:, :, :, 2]
output_content_features = tf.transpose(content_features)[2, :, :, :]
```

In [122]:

```
content_loss = feature_reconstruction_loss(content_img_features, output_con
tent_features)
```

In [123]:

```
total_style_loss = tf.zeros(1)
weight = 1.0 / len(style_layers)
for style_layer in style_layers:
    style_features = layers_dict[style_layer]
    # style_img_features = style_features[:, :, :, 1]
    # output_img_features = style_features[:, :, :, 2]
    style_img_features = tf.transpose(content_features)[1, :, :, :]
    output_img_features = tf.transpose(content_features)[2, :, :, :]
    total_style_loss += weight * style_loss(style_img_features,
output_img_features, (224, 224, 3))
```

Now we need to combine our two loss functions using the weightings we defined earlier. HINT: don't overthink this it should be a very simple operation.

In [124]:

```
total_loss = content_loss_weight*content_loss + style_loss_weight*total_style_loss
```

In [125]:

```
optimize = tf.train.AdamOptimizer(learning_rate=10).minimize(total_loss, var_list=[output_tensor])
```

Part 2: Feeding in Images

We now want to load and preprocess our images. keras provides a `load_img` function that conveniently loads our image and then cuts it down to our target size. Keras also provides a `vgg16.preprocess_input` that preprocesses images to be in the format vgg16 expects. Use these two functions to write the `load_image` function below.

In [126]:

```
from keras.preprocessing.image import load_img, img_to_array
from keras.applications.vgg16 import preprocess_input
import numpy as np

def load_image(img_path):
    img = load_img(target_size=(224,224,3),path=img_path) #YOUR CODE HERE
    call load_img and set the target size to be (224,224,3)
    img = img_to_array(img)
    img = np.expand_dims(img, axis=0)
    img = preprocess_input(img)
    return img

def deprocess_image(x):
    x = x.reshape((224, 224, 3))
    x[:, :, 0] += 103.939
    x[:, :, 1] += 116.779
    x[:, :, 2] += 123.68
    # 'BGR' -> 'RGB'
    x = x[:, :, ::-1]
    x = np.clip(x, 0, 255).astype('uint8')
    return x
```

In [127]:

```
content_img_path = 'images/sproul.jpg'
style_img_path = 'images/monet_style.jpg'
```

In [128]:

```
content_img = load_image(content_img_path)
style_img = load_image(style_img_path)
```

In [129]:


```
import matplotlib.pyplot as plt
%matplotlib inline
```

In [130]:

```
plt.imshow(deprocess_image(content_img))
```

Out[130]:

<matplotlib.image.AxesImage at 0x11c938080>

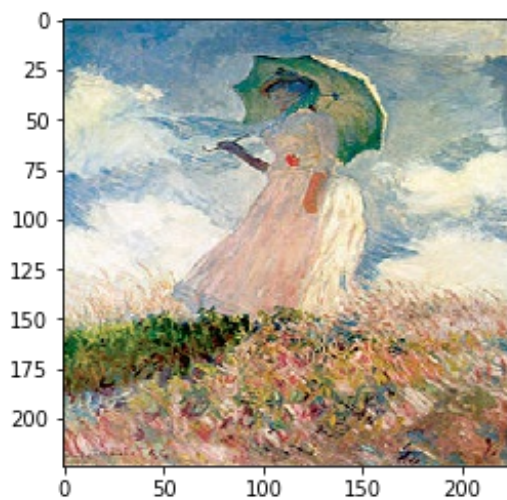


In [131]:

```
plt.imshow(deprocess_image(style_img))
```

Out[131]:

<matplotlib.image.AxesImage at 0x11c9cd3c8>



Part 3: Stylize Images

In [132]:

```
assign_var = tf.assign(output_tensor, content_img)
sess = K.get_session()
var = sess.run(assign_var)
```

Running the cell below will update the image 10 times. Since the initialization code is in the cell above, if you run the cell below and your output isn't great you can run for another 10 iterations simply by rerunning the cell below.

In [133]:

```
n_iterations = 10
```

```

for i in range(n_iterations):
    print("Running iteration: {}".format(i))
    _, output_val, loss = sess.run([optimize, output_tensor, total_loss], feed_dict={content_input: content_img, style_input: style_img})

```

```

Running iteration: 0
Running iteration: 1
Running iteration: 2
Running iteration: 3
Running iteration: 4
Running iteration: 5
Running iteration: 6
Running iteration: 7
Running iteration: 8
Running iteration: 9

```

In [134]:


```
output_img = deprocess_image(output_val)
```

In [135]:

```
plt.imshow(output_img)
```

Out[135]:

<matplotlib.image.AxesImage at 0x117cb6160>



Part 4: Style Transfer Writeup

Now you need to writeup your project. First, write a short paragraph about your understanding of how style transfer works. Feel free to refer to the paper if it helps but your paragraph needs to be in your own words.

Then attach 3 sets of images to your writeup. For each set show the original content image, the original style image, and the style transfer result. One set should be the images we provided here, include the content and style layers you used as well as the content and style weights you used. Another set should be the images we provided here but with different content layers, style layers and different content and style weights, include your choices for the layers and weights in your writeup. Finally, include a set of images that is based on a new content image and a new style image that you choose yourself. There will be an award for the group with the coolest style transfer result.

Set 1

Content Layer: 'block2_conv2'

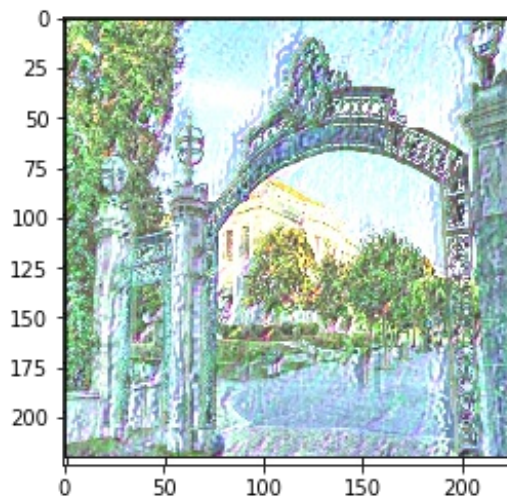
Three Styles: 'block2_pool', 'block3_conv1', 'block3_conv2'

In [137]:

```
plt.imshow(deprocess_image(load_image("images/satherGateStyle.png"))) #styl  
ized
```

Out [137]:

<matplotlib.image.AxesImage at 0x11bb24940>

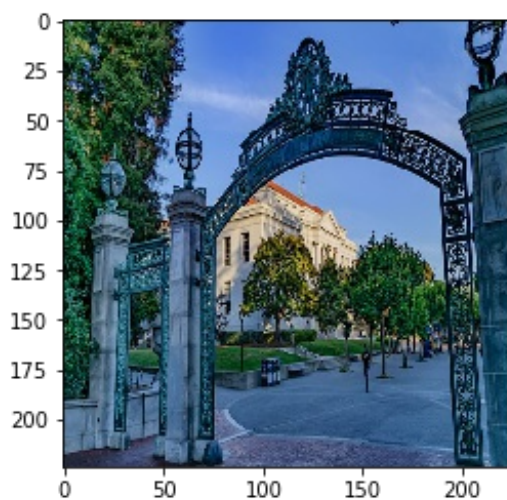


In [138]:

```
plt.imshow(deprocess_image(load_image("images/satherGate.jpg"))) #original
```

Out [138]:

<matplotlib.image.AxesImage at 0x11cb12d68>



Set 2

Content Layer: 'block1_conv2'

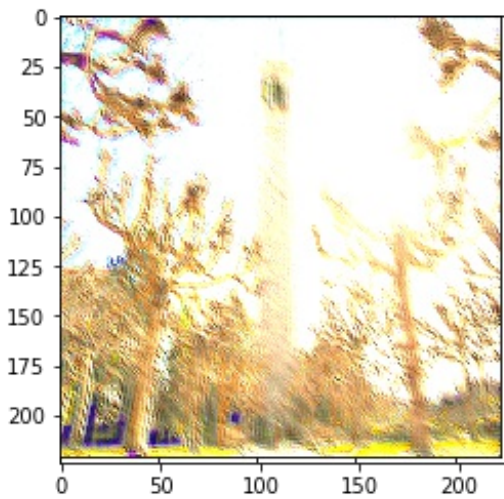
Three Layers: 'block1_pool', 'block2_conv1', 'block2_conv2'

In [112]:

```
plt.imshow(deprocess_image(load_image("images/campBerkStyle.png"))) #stylized
```

Out[112]:

<matplotlib.image.AxesImage at 0x11b381908>

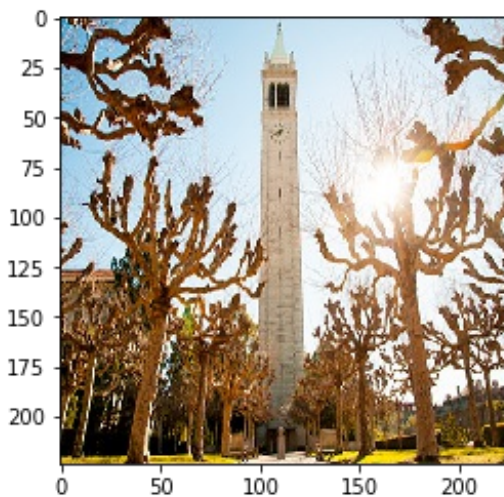


In [113]:

```
plt.imshow(deprocess_image(load_image("images/campBerk.jpg"))) #original
```

Out[113]:

<matplotlib.image.AxesImage at 0x11b3e1cc0>



Set 3

Content Layer: 'block2_conv1'

Three Layers: 'block5_conv2', 'block5_conv3', 'block5_pool'

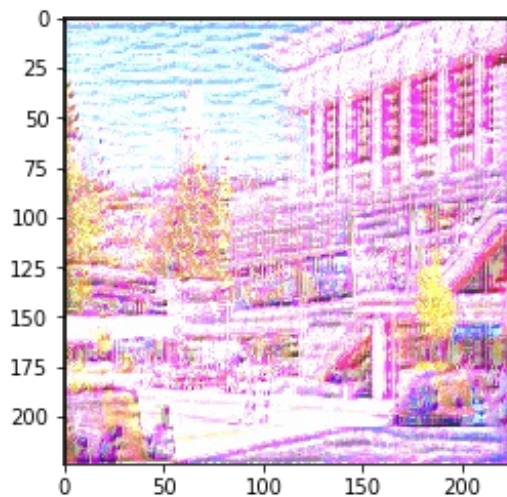
In [139]:

```
plt.imshow(deprocess_image(load_image("images/sproulStyle.png"))) #stylized
```

Out[139]:

<matplotlib.image.AxesImage at 0x11da53208>

<matplotlib.image.AxesImage at 0x1ba99280>




In [136]:

```
plt.imshow(deprocess_image(load_image("images/sproul.jpg"))) #original
```

Out [136]:

<matplotlib.image.AxesImage at 0x11ba49518>



Style transfer works by starting off with a random noise image, then after every iteration our original image changes a bit by bit to eventually adopt the style of the other image. Using convolution neural networks we use backpropagation to minimize our defined loss functions and then use the gradient of these loss function to compute our output image. We played around with different layer values to see the different results we can get.